

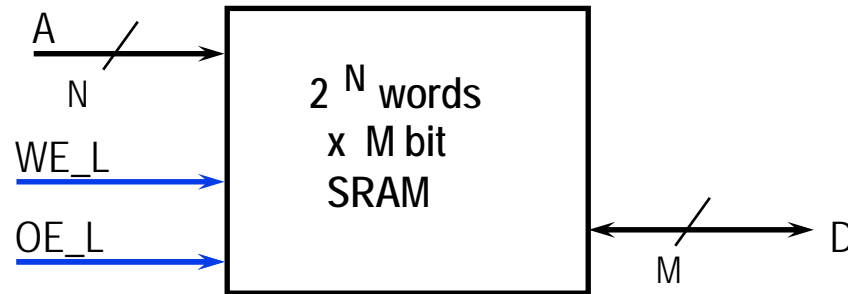
Computer Architecture

Ch. 7-2: Cache System

Spring, 2005

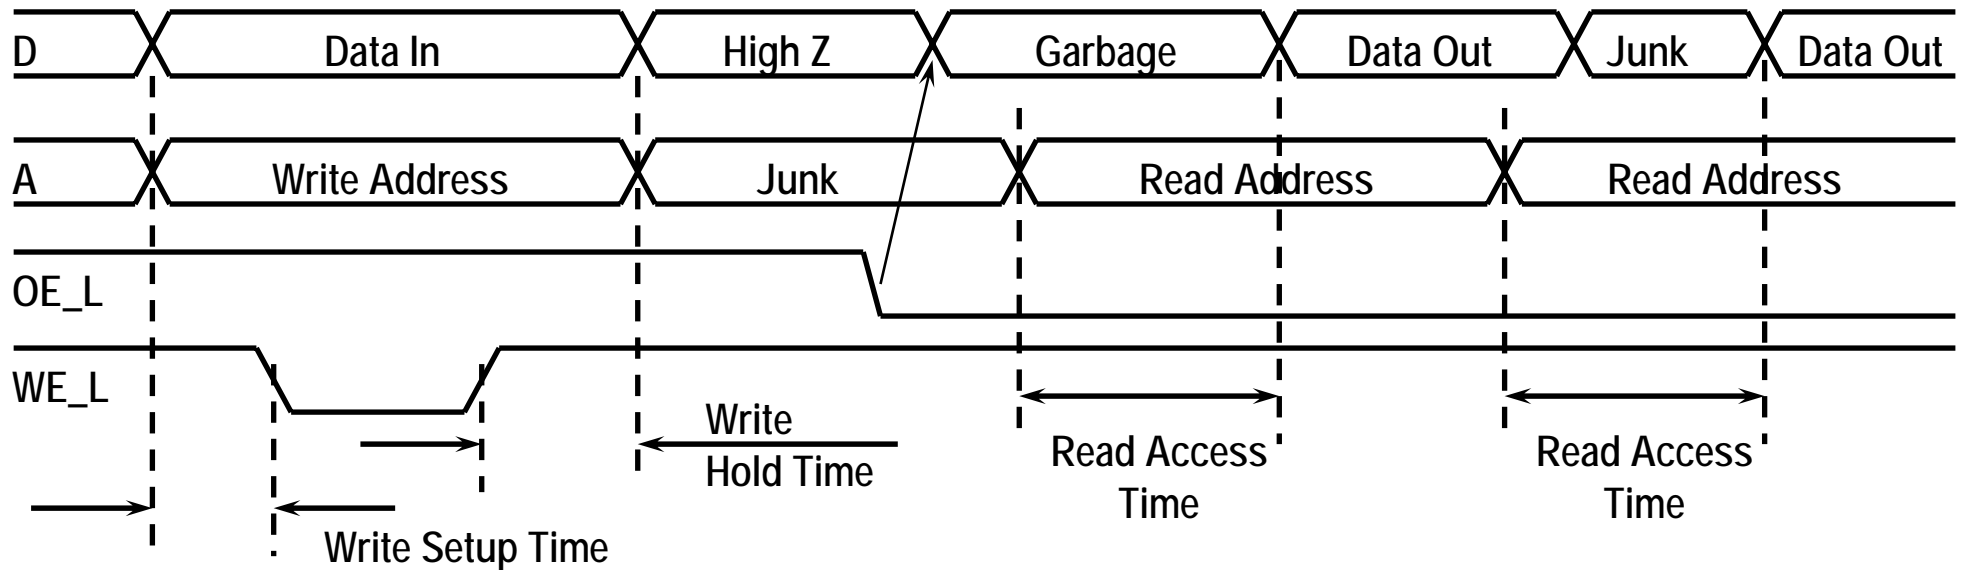
Sao-Jie Chen (csj@cc.ee.ntu.edu.tw)

Recap: SRAM Timing




Write Timing:

Read Timing:



The Motivation for Caches



- **Motivation:**
 - Large memories (DRAM) are slow
 - Small memories (SRAM) are fast
- Make the *average access time* small by:
 - Servicing most accesses from a small, fast memory.
- Reduce the *bandwidth* required of the large memory

Levels of the Memory Hierarchy

Capacity
Access Time
Cost

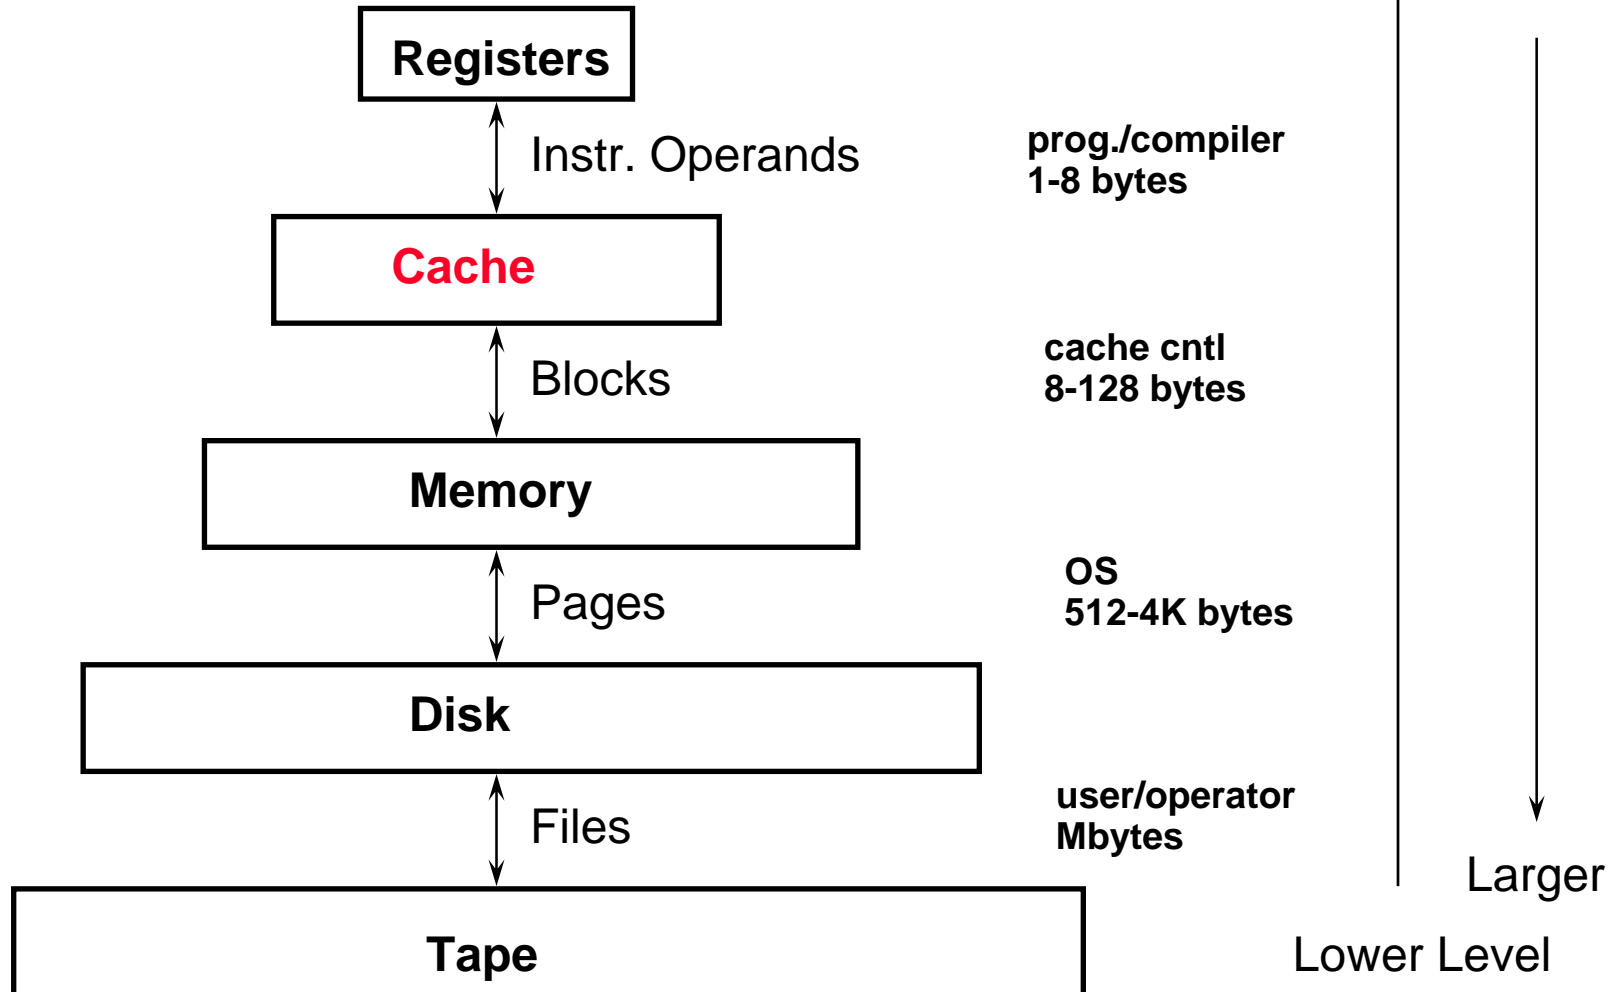
CPU Registers
100s Bytes
<10s ns

Cache
K Bytes
10-100 ns
\$.01-.001/bit

Main Memory
M Bytes
100ns-1us
\$.01-.001

Disk
G Bytes
ms₋₃ - 10⁻⁴
10⁻³ - 10 cents

Tape
infinite
sec-min
10⁻⁶



Locality


- A principle that makes having a memory hierarchy a good idea
- If an item is referenced,
 - temporal locality:** it will tend to be referenced again soon
 - spatial locality:** nearby items will tend to be referenced soon.

Why does code have locality?

- Our initial focus: two levels (upper, lower)
 - block: minimum unit of data
 - hit: data requested is in the upper level
 - miss: data requested is not in the upper level

Memory Hierarchy: Terminology

- **Hit**: data appears in some block in the upper level (example: Block X)
 - **Hit Rate**: the fraction of memory access found in the upper level
 - **Hit Time**: Time to access the upper level which consists of
RAM access time + Time to determine hit/miss
- **Miss**: data needs to be retrieved from a block in the lower level (Block Y)
 - **Miss Rate** = $1 - (\text{Hit Rate})$
 - **Miss Penalty** = Time to replace a block in the upper level +
Time to deliver the block to the processor
- **Hit Time** \ll **Miss Penalty**




Basic Terminology: Typical Values

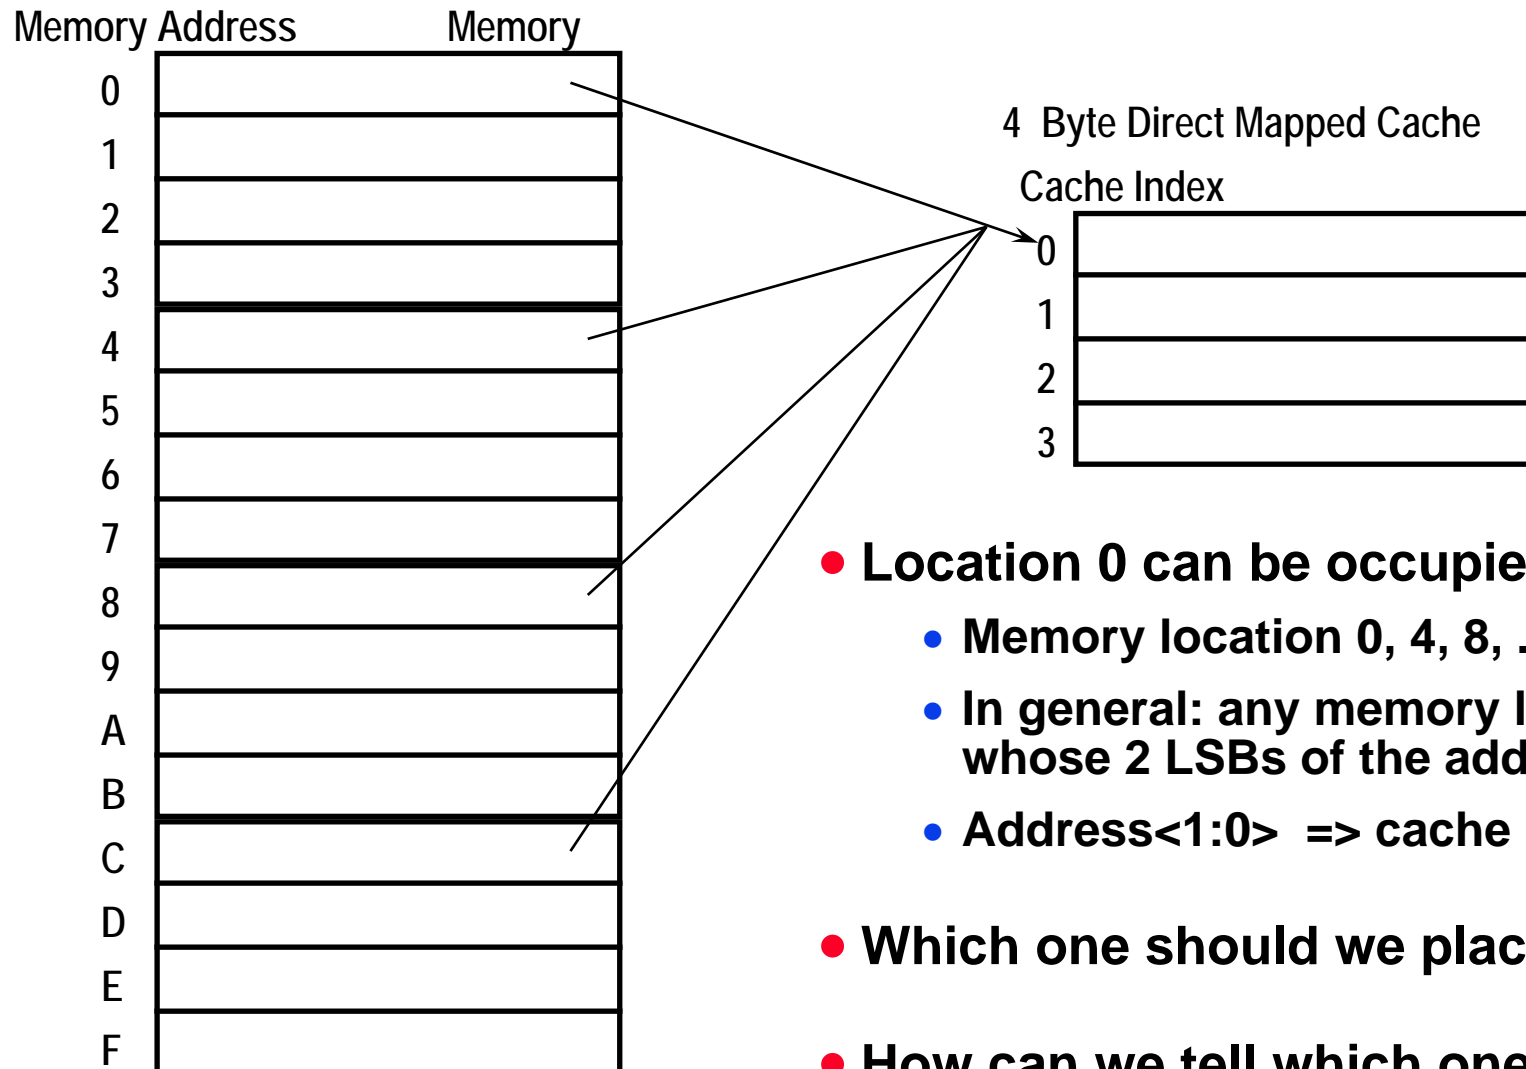
Typical Values	
Block (line) size	4 - 128 bytes
Hit time	1 - 4 cycles
Miss penalty	8 - 32 cycles (and increasing)
(access time)	(6-10 cycles)
(transfer time)	(2 - 22 cycles)
Miss rate	1% - 20%
Cache Size	1 KB - 256 KB

How Does Cache Work?

- **Temporal Locality** (Locality in Time): If an item is referenced, it will tend to be referenced again soon.
 - Keep more recently accessed data items closer to the processor
- **Spatial Locality** (Locality in Space): If an item is referenced, items whose addresses are close by tend to be referenced soon.
 - Move blocks consists of contiguous words to the cache



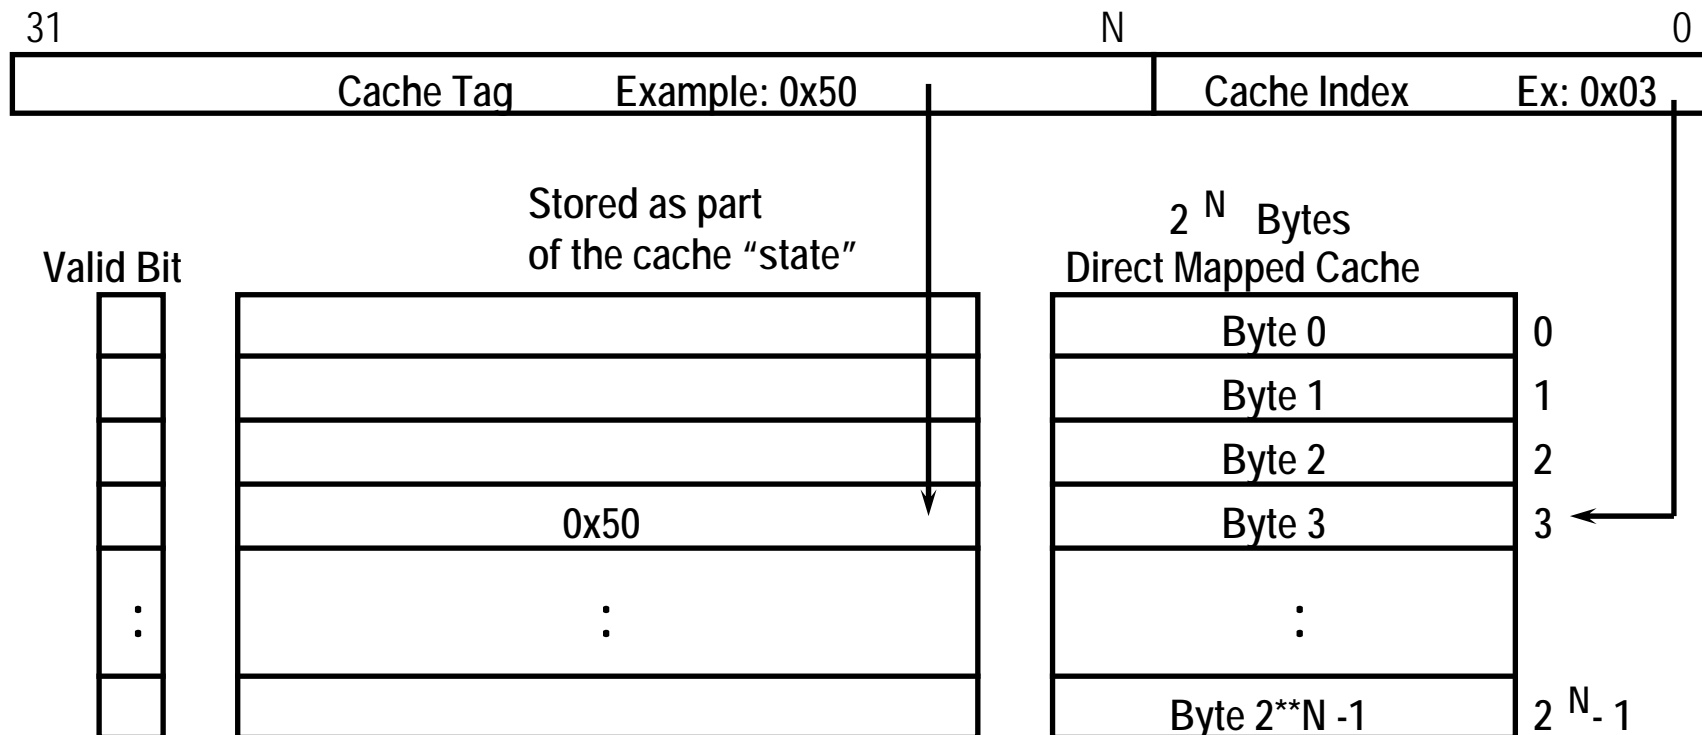
The Simplest Cache: Direct Mapped Cache






- **Location 0 can be occupied by data from:**
 - Memory location 0, 4, 8, ... etc.
 - In general: any memory location whose 2 LSBs of the address are 0s
 - Address<1:0> => cache index
- **Which one should we place in the cache?**
- **How can we tell which one is in the cache?**

Cache Tag and Cache Index

- Assume a 32-bit memory (byte) address:
 - A 2^N bytes direct mapped cache:
 - ⇒ Cache Index: The lower N bits of the memory address
 - ⇒ Cache Tag: The upper $(32 - N)$ bits of the memory address



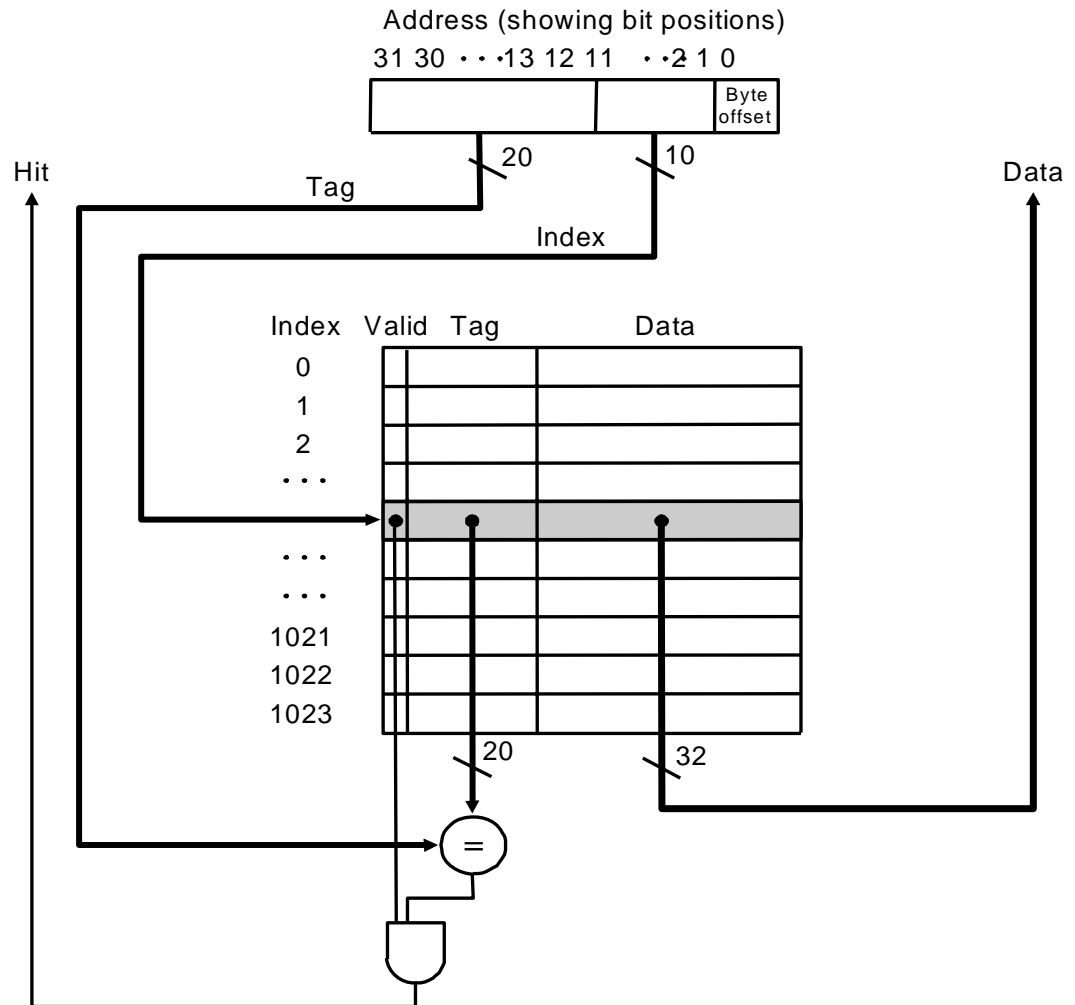
Cache Access Example



- **Sad Fact of Life:**
 - A lot of misses at start up:
Compulsory Misses
⇒ **(Cold start misses)**

Direct Mapped Cache


- For MIPS:



What kind of locality are we taking advantage of?




Definition of a Cache Block

- **Cache Block**: the cache data that has in its own cache tag
- Our previous “extreme” example:
 - 4-byte Direct Mapped cache: Block Size = 1 Byte
 - Take advantage of Temporal Locality: If a byte is referenced, it will tend to be referenced soon.
 - Did not take advantage of Spatial Locality: If a byte is referenced, its adjacent bytes will be referenced soon.
- In order to take advantage of Spatial Locality: increase the block size



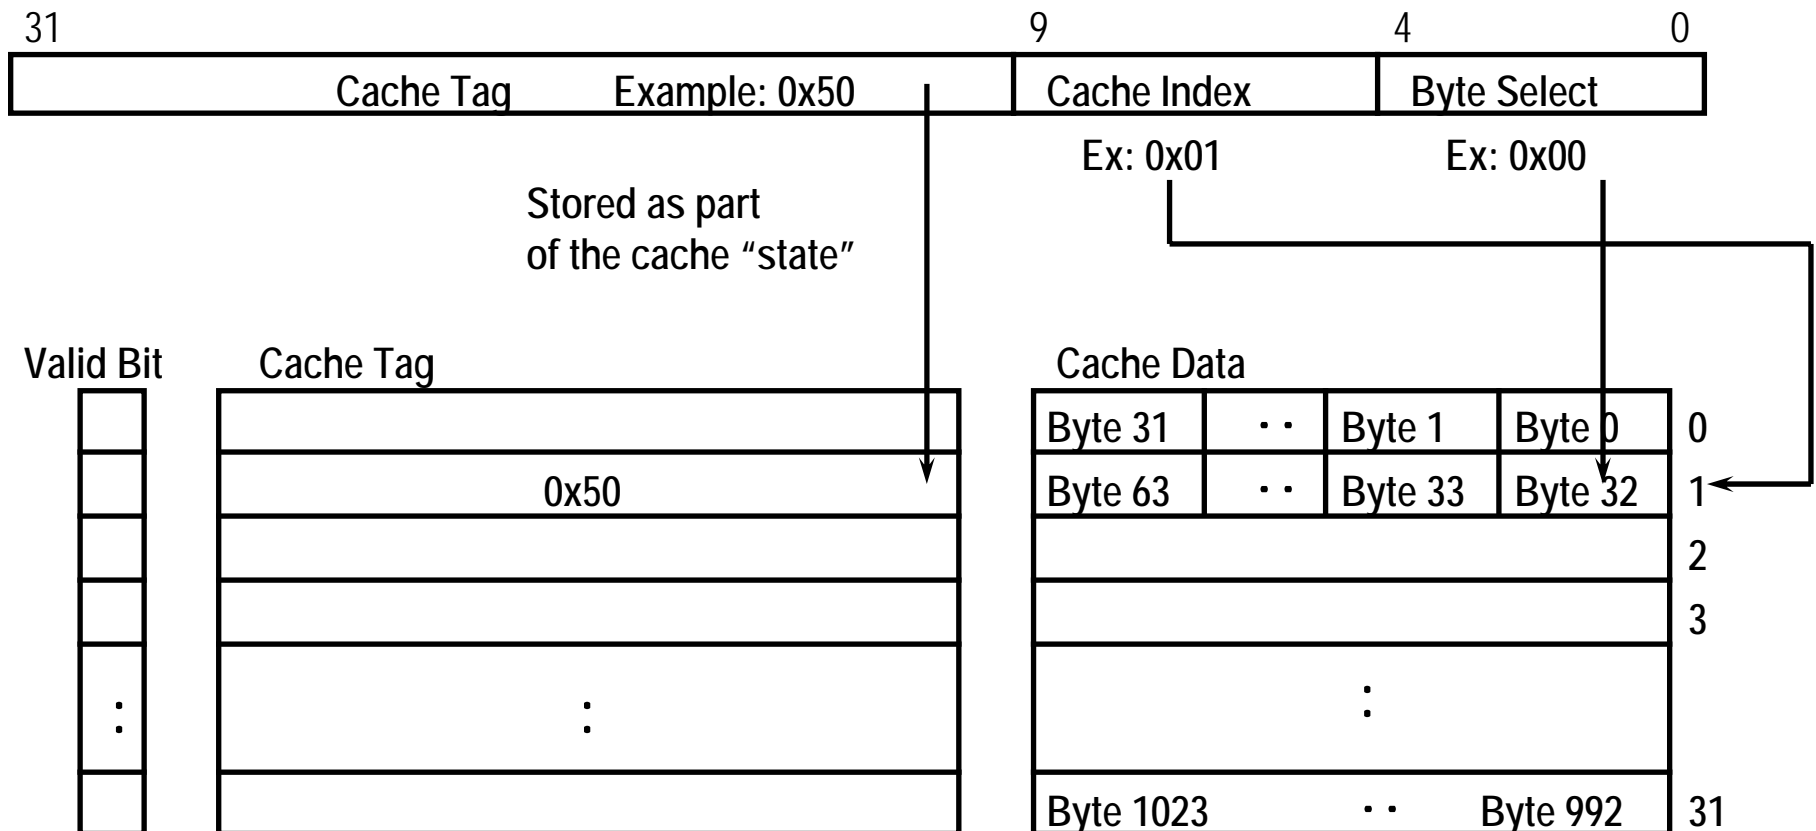
Block Size Tradeoff

- In general, larger block size take advantage of spatial locality BUT:
 - Larger block size means larger miss penalty:
 - ⇒ Takes longer time to fill up the block
 - If block size is too big relative to cache size, miss rate will go up
- **Average Access Time:**
$$= \text{Hit Time} \times (1 - \text{Miss Rate}) + \text{Miss Penalty} \times \text{Miss Rate}$$



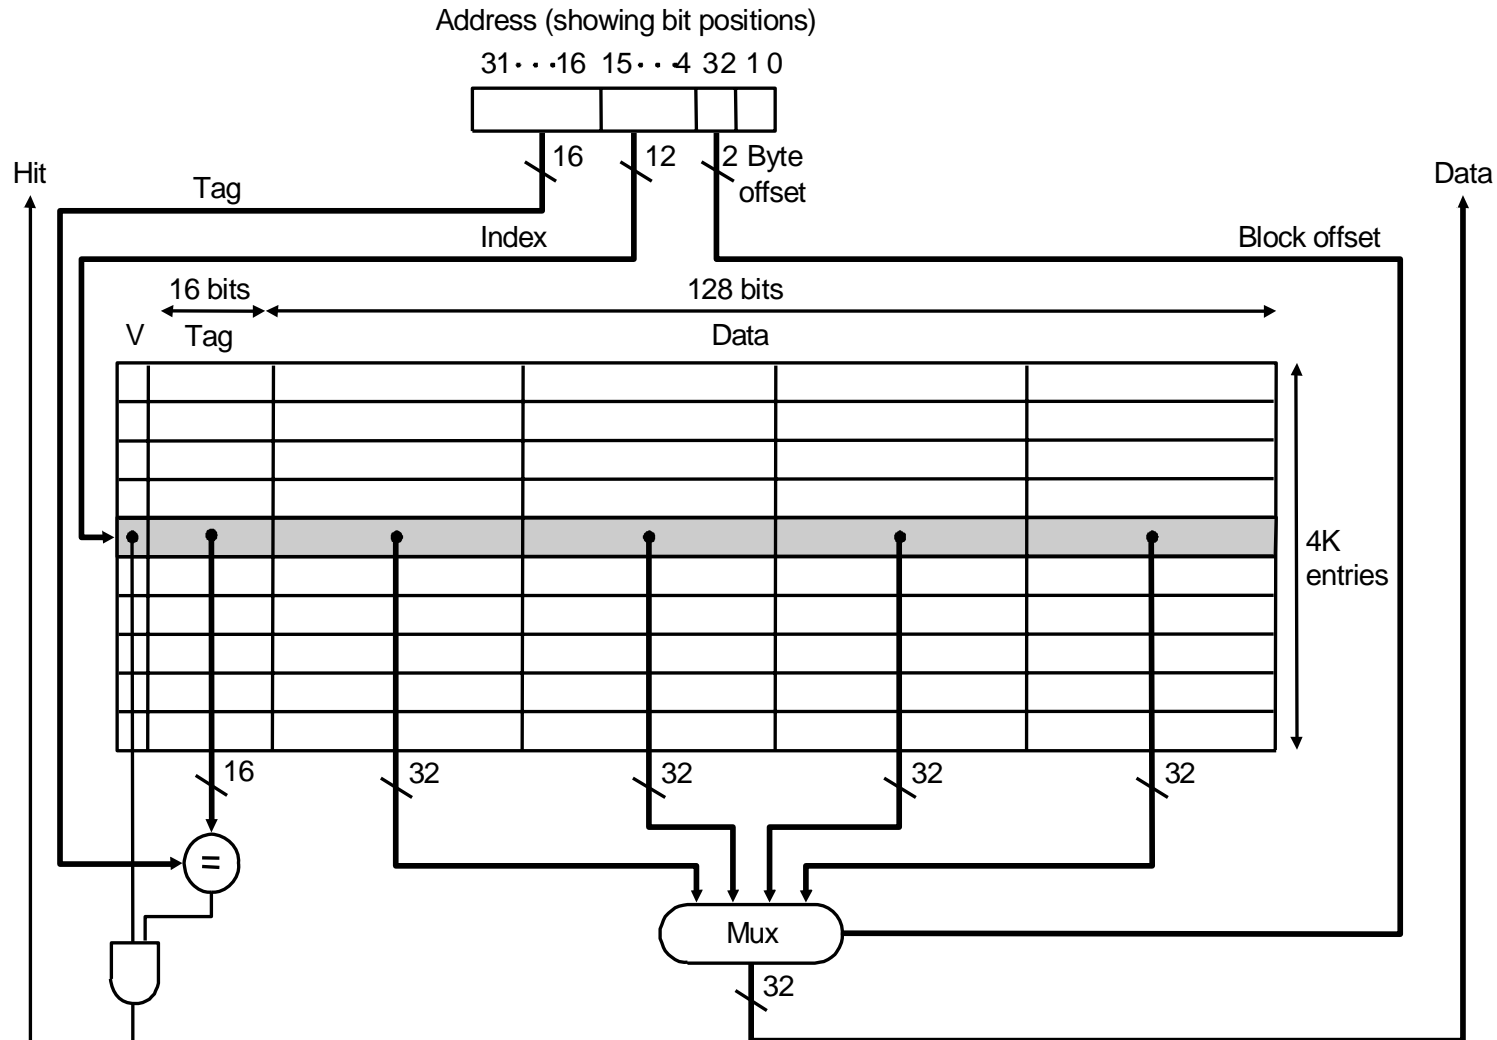
Example: 1 KB Direct Mapped Cache with 32 B Blocks

- For a 2^N byte cache:
 - The uppermost $(32 - N)$ bits are always the Cache Tag
 - The lowest M bits are the Byte Select (Block Size = 2^M)




Direct Mapped Cache

- Taking advantage of spatial locality:



Performance

- Increasing the block size tends to decrease miss rate:



- Use split caches because there is more spatial locality in code:

Program	Block size in words	Instruction miss rate	Data miss rate	Effective combined miss rate
gcc	1	6.1%	2.1%	5.4%
	4	2.0%	1.7%	1.9%
spice	1	1.2%	1.3%	1.2%
	4	0.3%	0.6%	0.4%

Performance

- Simplified model:

$\text{execution time} = (\text{execution cycles} + \text{stall cycles}) \times \text{cycle time}$


$\text{stall cycles} = \# \text{ of instructions} \times \text{miss ratio} \times \text{miss penalty}$

- Two ways of improving performance:

- decreasing the miss ratio (e.g., use Associativity)
- decreasing the miss penalty (e.g., use Multi-level cache)

What happens if we increase block size?

Another Extreme Example



- **Cache Size = 4 bytes** **Block Size = 4 bytes**
 - Only ONE entry in the cache
- **True: If an item is accessed, likely that it will be accessed again soon**
 - But it is unlikely that it will be accessed again immediately!!!
 - The next access will likely to be a miss again
 - ⇒ Continually loading data into the cache but discard (force out) them before they are used again
 - ⇒ Worst nightmare of a cache designer: Ping Pong Effect
- **Conflict Misses are misses caused by:**
 - Different memory locations mapped to the same cache index
 - ⇒ Solution 1: **Make the cache size bigger**
 - ⇒ Solution 2: **Multiple entries for the same Cache Index**

Decreasing miss ratio with associativity

One-way set associative
(direct mapped)

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Two-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

Four-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

Eight-way set associative (fully associative)


Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data

- ***Compared to direct mapped, give a series of references that:***
 - ***results in a lower miss ratio using a 2-way set associative cache***
 - ***results in a higher miss ratio using a 2-way set associative cache***

assuming we use the “least recently used” replacement strategy


A Two-way Set Associative Cache

- **N-way set associative:** N entries for each Cache Index
 - N direct mapped caches operates in parallel
- **Example: Two-way set associative cache**
 - Cache Index selects a “set” from the cache
 - The two tags in the set are compared in parallel
 - Data is selected based on the tag result

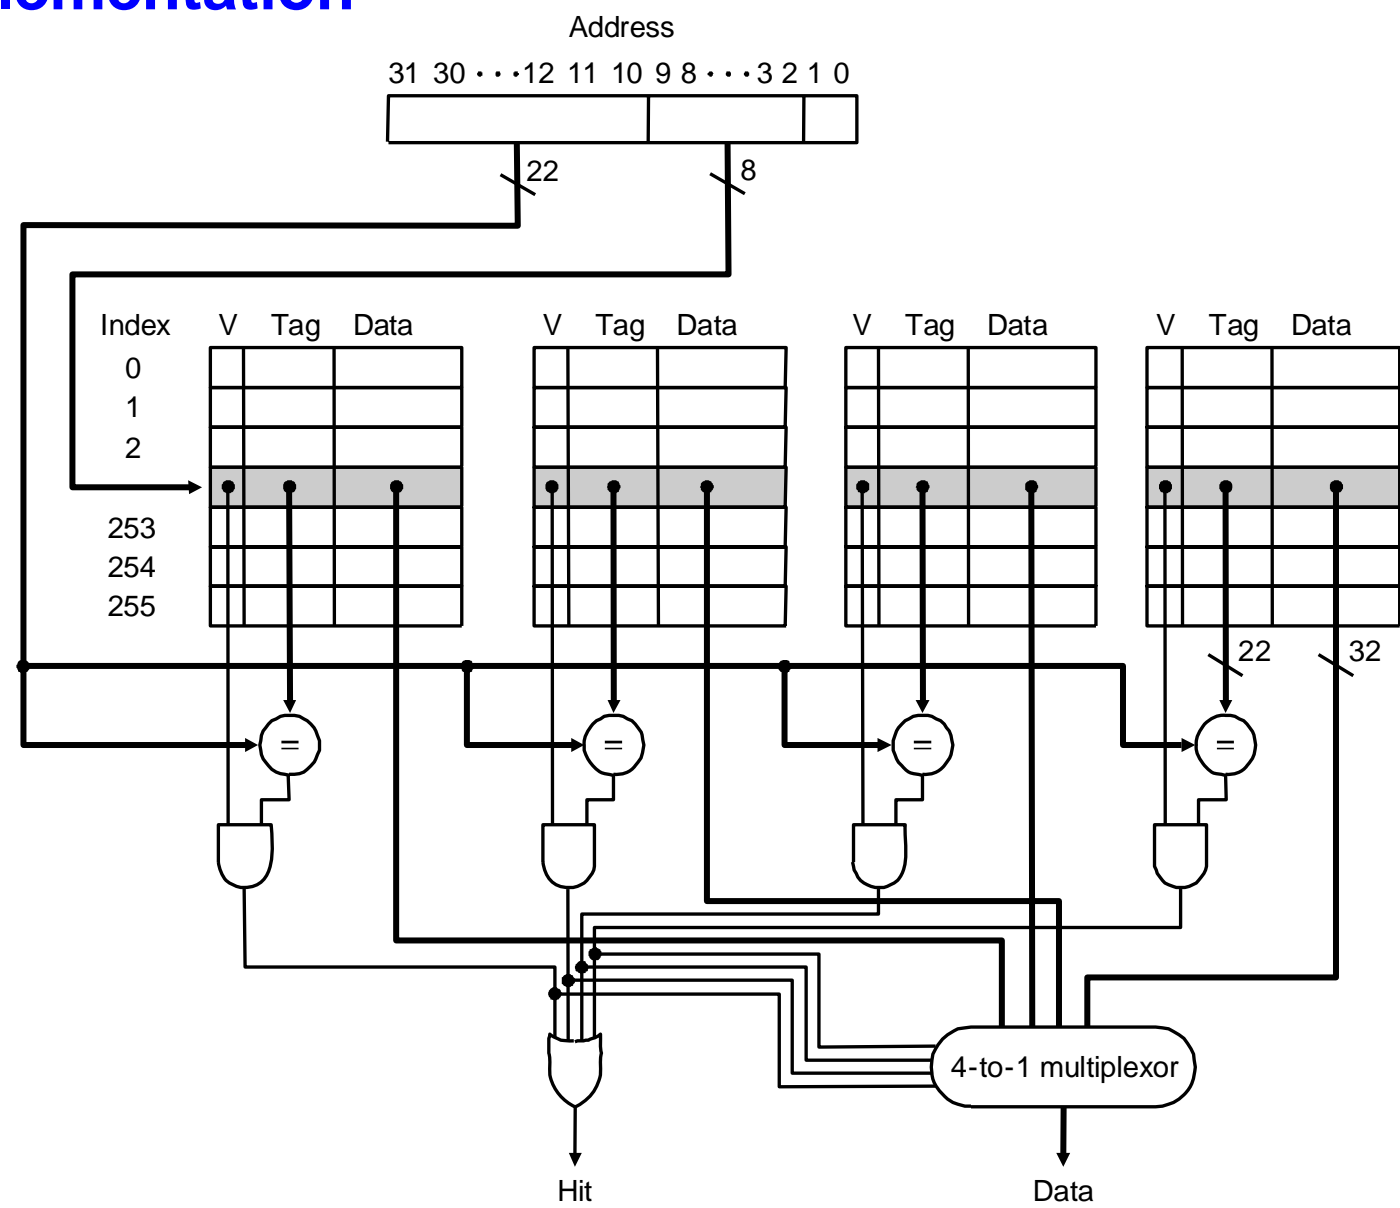


Disadvantage of Set Associative Cache


- **N-way Set Associative Cache** versus **Direct Mapped Cache**:
 - N comparators vs. 1
 - Extra MUX delay for the data
 - Data comes AFTER Hit/Miss
- In a direct mapped cache, Cache Block is available BEFORE Hit/Miss:
 - Possible to assume a hit and continue. Recover later if miss.



An implementation

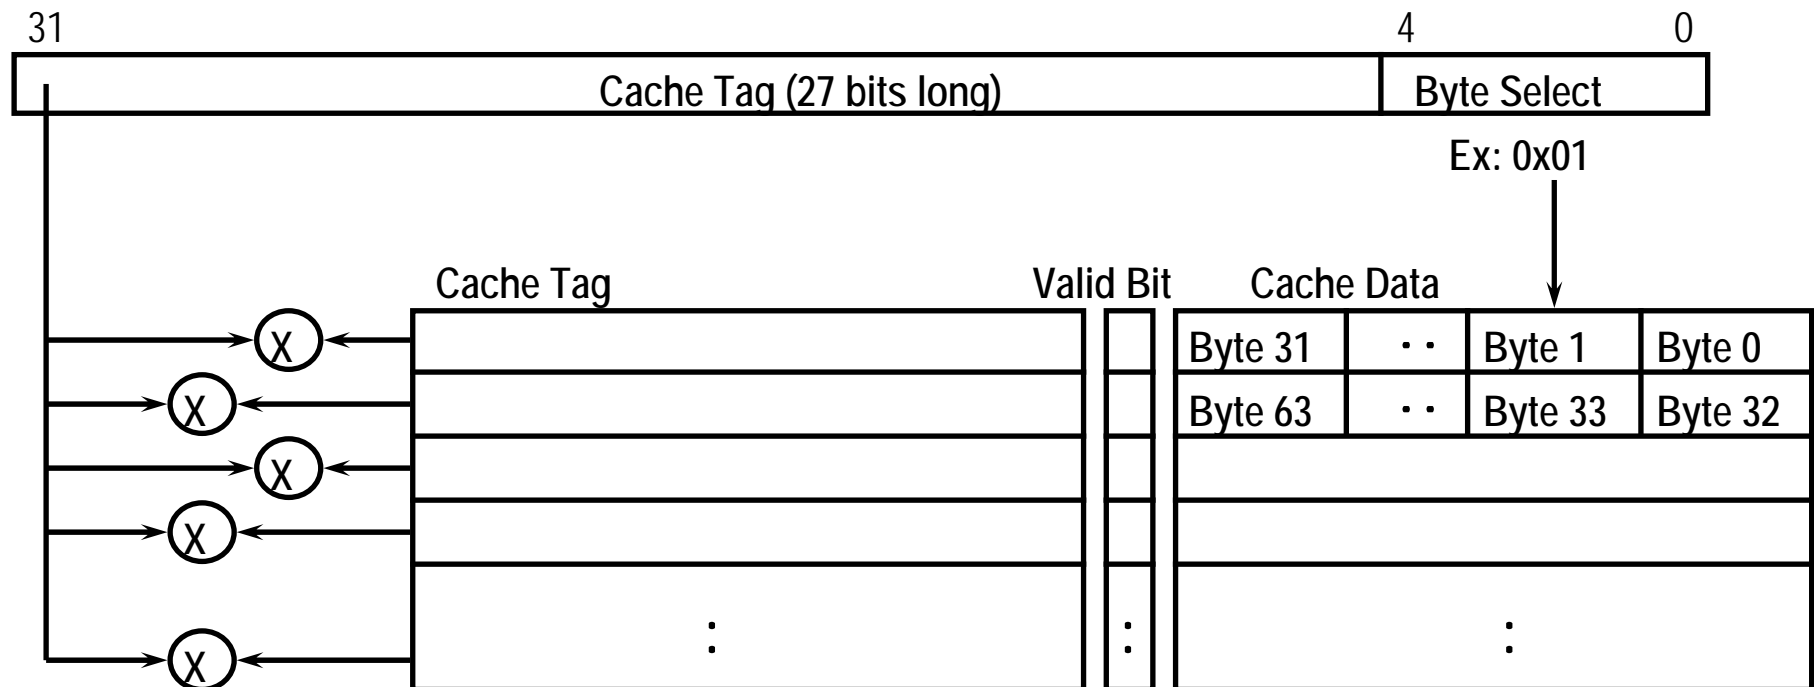


Performance



And yet Another Extreme Example: Fully Associative

- **Fully Associative Cache** -- push the set associative idea to its limit!
 - Forget about the Cache Index
 - Compare the Cache Tags of all cache entries in parallel
 - Example: Block Size = 32 B blocks, we need 2^N 27-bit comparators
- By definition: Conflict Miss = 0 for a fully associative cache



A Summary on Sources of Cache Misses

- **Compulsory** (cold start, first reference): first access to a block
 - “cold” fact of life: not a whole lot you can do about it
- **Conflict** (collision):
 - Multiple memory locations mapped to the same cache location
 - Solution 1: increase cache size
 - Solution 2: increase associativity
- **Capacity**:
 - Cache cannot contain all blocks access by the program
 - Solution: increase cache size
- **Invalidation**: other process (e.g., I/O) updates memory

The Need to Make a Decision!

- **Direct Mapped Cache:**
 - Each memory location can only mapped to 1 cache location
 - No need to make any decision :-)
 - ⇒ Current item replaced the previous item in that cache location
- **N-way Set Associative Cache:**
 - Each memory location have a choice of N cache locations
- **Fully Associative Cache:**
 - Each memory location can be placed in ANY cache location
- **Cache miss in a N-way Set Associative or Fully Associative Cache:**
 - Bring in new block from memory
 - Throw out a cache block to make room for the new block
 - Damn! We need to make a decision on which block to throw out!

Cache Block Replacement Policy

- **Random Replacement:**


- Hardware randomly selects a cache item and throw it out

- **Least Recently Used:**

- Hardware keeps track of the access history
- Replace the entry that has not been used for the longest time

- **Example of a Simple “Pseudo” Least Recently Used Implementation:**


- Assume 64 Fully Associative Entries
- Hardware replacement pointer points to one cache entry
- Whenever an access is made to the entry the pointer points to:
 - ⇒ Move the pointer to the next entry
- Otherwise: do not move the pointer



Cache Write Policy: Write Through versus Write Back


- Cache read is much easier to handle than cache write:
 - Instruction cache is much easier to design than data cache
- Cache write:
 - How do we keep data in the cache and memory consistent?
- Two options (decision time again :-)
 - **Write Back**: write to cache only. Write the cache block to memory when that cache block is being replaced on a cache miss.
 - ⇒ Need a “dirty” bit for each cache block
 - ⇒ Greatly reduce the memory bandwidth requirement
 - ⇒ Control can be complex
 - **Write Through**: write to cache and memory at the same time.
 - ⇒ What!!! How can this be? Isn't memory too slow for this?

Write Buffer for Write Through




- A **Write Buffer** is needed between the Cache and Memory
 - Processor: writes data into the cache and the write buffer
 - Memory controller: write contents of the buffer to memory
- Write buffer is just a FIFO:
 - Typical number of entries: 4
 - Works fine if: Store frequency (w.r.t. time) $\ll 1 / \text{DRAM write cycle}$
- Memory system designer's nightmare:
 - Store frequency (w.r.t. time) $\gg 1 / \text{DRAM write cycle}$
 - Write buffer saturation

Write Buffer Saturation



- **Store frequency (w.r.t. time) $\gg 1 / \text{DRAM write cycle}$**
 - If this condition exists for a long period of time (CPU cycle time too quick and/or too many store instructions in a row):
 - \Rightarrow Store buffer will overflow no matter how big you make it
 - \Rightarrow The CPU Cycle Time \leq DRAM Write Cycle Time
- **Solution for write buffer saturation:**
 - Use a write back cache
 - Install a second level (L2) cache:



Decreasing miss penalty with multilevel caches

- **Add a second level cache:**
 - often primary cache is on the same chip as the processor
 - use SRAMs to add another cache above primary memory (DRAM)
 - miss penalty goes down if data is in 2nd level cache
- **Example:**
 - CPI of 1.0 on a 500Mhz machine with a 5% miss rate, 200ns DRAM access
 - Adding 2nd level cache with 20ns access time decreases miss rate to 2%
- **Using multilevel caches:**
 - try and optimize the hit time on the 1st level cache
 - try and optimize the miss rate on the 2nd level cache

Decreasing miss penalty with multilevel caches

- **Example:**


- CPI of 1.0 on a 500Mhz machine with a 5% miss rate, 200ns DRAM access
- Adding 2nd level cache with 20ns access time decreases miss rate to 2%

- **Solution:**


- 500MHz = 2 ns /clock cycle (cc);
- Miss penalty to main memory is: $200 \text{ ns} / (2\text{ns/cc}) = 100 \text{ cc}$
- Miss penalty to 2nd level cache: $20 \text{ ns} / (2\text{ns/cc}) = 10 \text{ cc}$
- One-level cache: $\text{CPI} = 1.0 + 5\% \times 100 = 6.0$
- Two-level cache: $\text{CPI} = 1.0 + 5\% \times 10 + 2\% \times 100 = 3.5$
- Two-level cahce is $6.0 / 3.5 = 1.7$ times faster.

Cache Complexity

- Not always easy to understand implications of caches:



**Theoretical behavior of
Radix sort vs. Quicksort**



**Observed behavior of
Radix sort vs. Quicksort**

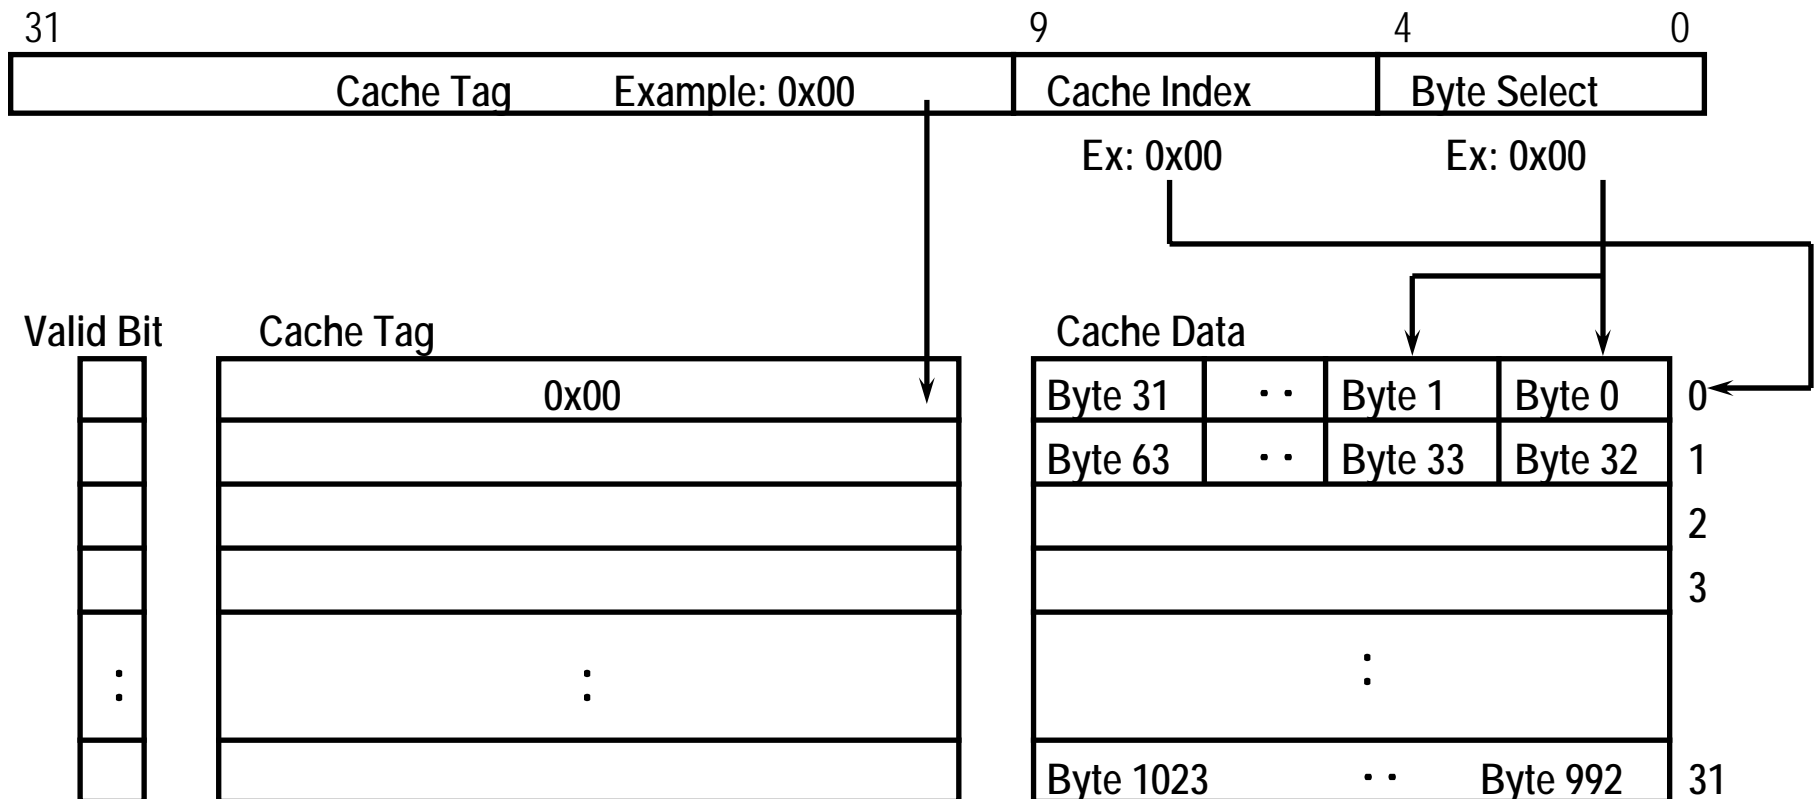
Write Allocate versus Not Allocate

- Assume: a 16-bit write to memory location 0x0 and causes a miss

- Do we read in the rest of the block (Byte 2, 3, ... 31)?

Yes: Write Allocate

No: Write Not Allocate




What is a Sub-block?

- **Sub-block:**


- A unit within a block that has its own valid bit
- Example: 1 KB Direct Mapped Cache, 32-B Block, 8-B Sub-block

⇒ Each cache entry will have: $32/8 = 4$ valid bits


- Write miss: only the bytes in that sub-block is brought in.




Reducing Memory Transfer Time



Solution 1
High BW DRAM



Solution 2
Wide Path Between Memory & Cache



Solution 3
Memory Interleaving

Examples:

Page Mode DRAM
SDRAM
CDRAM
RAMbus

Cost
↓

Summary:

- **The Principle of Locality:**

- Program access a relatively small portion of the address space at any instant of time.

- ⇒ **Temporal Locality**: Locality in Time

- ⇒ **Spatial Locality**: Locality in Space

- **Three Major Categories of Cache Misses:**

- **Compulsory Misses**: sad facts of life. Example: cold start misses.
- **Conflict Misses**: increase cache size and/or associativity.
Nightmare Scenario: ping pong effect!
- **Capacity Misses**: increase cache size

- **Write Policy:**

- **Write Through**: need a write buffer. Nightmare: WB saturation
- **Write Back**: control can be complex

Where to get more information?

- **General reference, Chapter 8 of:**
 - John Hennessy & David Patterson, *Computer Architecture: A Quantitative Approach*, 3rd Ed., Morgan Kaufmann Publishers Inc., 2002
- **A landmark paper on caches:**
 - Alan Smith, “Cache Memories,” *Computing Surveys*, September 1982
- **A book on everything you need to know about caches:**
 - Steve Przybylski, *Cache and Memory Hierarchy Design: A Performance-Directed Approach*, Morgan Kaufmann Publishers Inc., 1990.